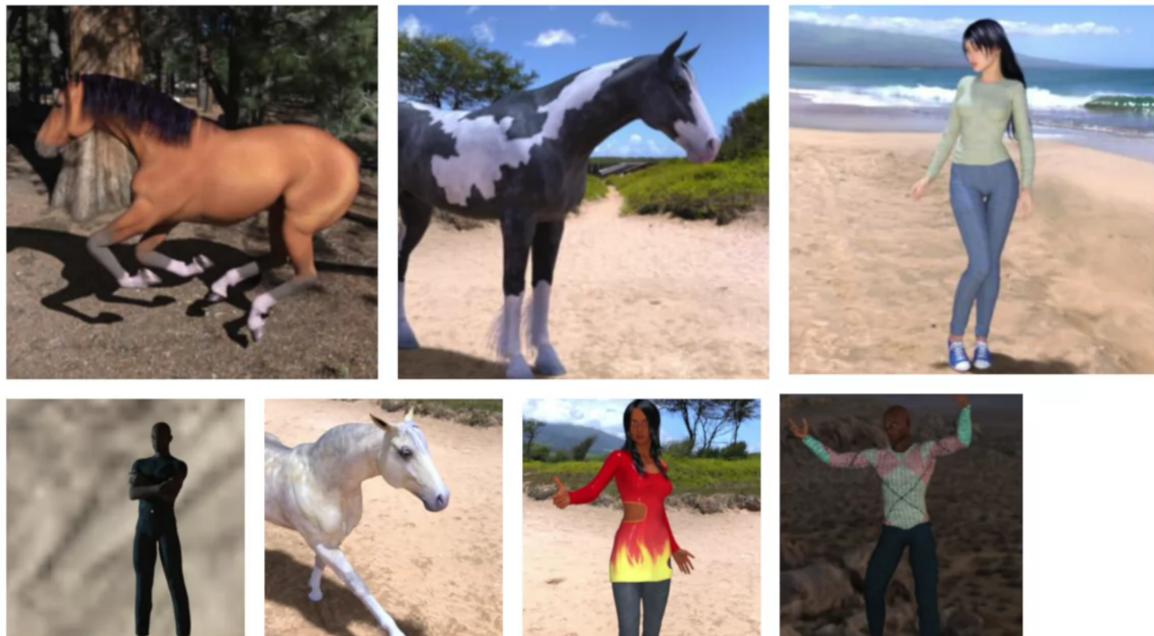




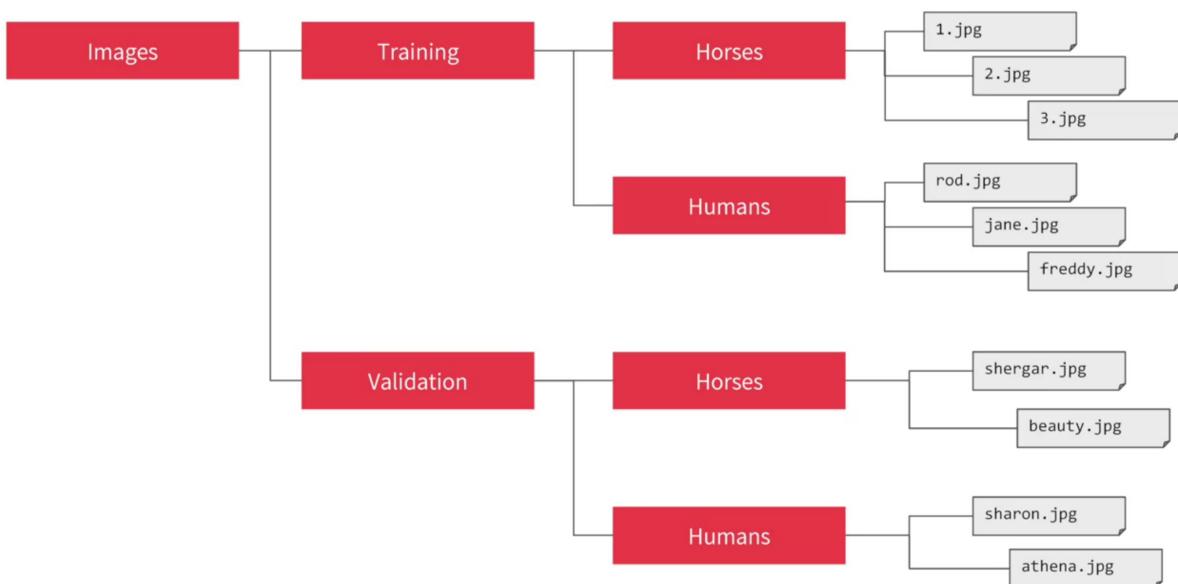
# Week 4

## Understanding Image Generator

- Previous image samples are all the same size and at the same location
- These images we will be using would be in different aspect ratios, different location and different sizes and even might have multiple subjects



- all you have to do to generate image is to point it at a directory



## Defining a ConvNet to use complex images

- The layers of the algorithm

```

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu',
                         input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

```

- summary of the method above

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d_5 (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_6 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_6 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_7 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten_1 (Flatten)	(None, 78400)	0
dense_2 (Dense)	(None, 512)	40141312
dense_3 (Dense)	(None, 1)	513
Total params:	40,165,409	
Trainable params:	40,165,409	
Non-trainable params:	0	

- Compiling the model

```

model.compile(loss='binary_crossentropy',
              optimizer=RMSprop(lr=0.001),
              metrics=['acc'])

```

- because dealing with binary here, so this is using binary crossentropy instead
- RMSprop adjusts the learning rate
- Fitting generator

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=8,  
    epochs=15,  
    validation_data=validation_generator,  
    validation_steps=8,  
    verbose=2)
```

- using fit\_generator because data using is generated with generator instead of dataset
- First parameter is the training generator set up earlier
- In order to load all the data in we need to use 8 epochs
- Epochs to train for is 15
- validation set is using the validation generator created earlier
- 8 steps so handle it by 32
- verbose specifies how much to display while training

```

import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():

    # predicting images
    path = '/content/' + fn
    img = image.load_img(path, target_size=(300, 300))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a human")
    else:
        print(fn + " is a horse")

```

- The code above is using the trained data
- Some specific for Google Colab

## Horse and Human identifier

```

import tensorflow as tf
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import os

class MyCallBack(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, log):
        if log.get('accuracy') > 0.999:
            print("\nReached 99.9 accuracy so cancelling training!")
            self.model.stop_training = True

callback = MyCallBack()

train_hourse_dir = os.path.join('C:/Users/Alex Lai.DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-
train_human_dir = os.path.join('C:/Users/Alex Lai.DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-0

train_hourse_labels = os.listdir(train_hourse_dir)
train_human_labels = os.listdir(train_human_dir)
# print(train_hourse_labels[0:10])
# print(train_human_labels[0:10])
# print('Total hourse Images', len(train_hourse_labels))
# print('Total human Images', len(train_human_labels))

# -----
# Plotting plt images
# nrows = 4
# ncols = 4
# pic_index = 0
# fig = plt.gcf()
# fig.set_size_inches(ncols*4, nrows*4)
# pic_index += 28
# next_horse_pix = [os.path.join(train_hourse_dir, fname)
#                   for fname in train_hourse_labels[pic_index-8:pic_index]]
# next_human_pix = [os.path.join(train_human_dir, fname)
#                   for fname in train_human_labels[pic_index-8:pic_index]]

```

```

# for i, img_path in enumerate(next_horse_pix+next_human_pix):
#     # Set up subplot; subplot indices start at 1
#     sp = plt.subplot(nrows, ncols, i + 1)
#     sp.axis('Off') # Don't show axes (or gridlines)

#     img = mpimg.imread(img_path)
#     plt.imshow(img)

# plt.show()
# ----- #

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(loss='binary_crossentropy', optimizer=RMSprop(lr=0.001), metrics=['accuracy'])
train_datagen = ImageDataGenerator(rescale=1/255)
train_generator = train_datagen.flow_from_directory(
    'C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-or-human/',
    target_size=(300, 300), batch_size=128, class_mode='binary')
history = model.fit(train_generator, steps_per_epoch=8, epochs=15, verbose=1, callbacks=[callback])

```

## Horse and Human identifier with validation

```

import tensorflow as tf
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import os

class MyCallBack(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, log):
        if log.get('accuracy') > 0.999:
            print("\nReached 99.9 accuracy so cancelling training!")
            self.model.stop_training = True

callback = MyCallBack()

train_hourse_dir = os.path.join('C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-train')
train_human_dir = os.path.join('C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-o')
val_hourse_dir = os.path.join('C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/validati')
val_human_dir = os.path.join('C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/validatio')

train_hourse_labels = os.listdir(train_hourse_dir)
train_human_labels = os.listdir(train_human_dir)
val_hourse_labels = os.listdir(val_hourse_dir)
val_human_labels = os.listdir(val_human_dir)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPool2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(loss='binary_crossentropy', optimizer=RMSprop(lr=0.001), metrics=['accuracy'])
train_datagen = ImageDataGenerator(rescale=1/255)
train_generator = train_datagen.flow_from_directory(
    'C:/Users/Alex Lai DESKTOP-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/horse-or-human/',

```

```

    target_size=(300, 300), batch_size=128, class_mode='binary')
val_generator = train_datagen.flow_from_directory(
    'C:/Users/Alex Lai/Desktop-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/horse or human/validation-horse-or-human/',
    target_size=(300, 300), batch_size=32, class_mode='binary')
history = model.fit(train_generator, steps_per_epoch=8, epochs=15, verbose=1, callbacks=[callback], validation_data=val_generator)

```

## Weekly test

```

import tensorflow as tf
import os
from os import path, getcwd, chdir
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import RMSprop

# GRADED FUNCTION: train_happy_sad_model
def train_happy_sad_model():
    # Please write your code only where you are indicated.
    # please do not remove # model fitting inline comments.

    DESIRED_ACCURACY = 0.999

    class myCallback(tf.keras.callbacks.Callback):
        def on_epoch_end(self, epoch, log={}):
            if log.get('accuracy') > DESIRED_ACCURACY:
                print('Reached 99.9% accuracy so cancelling training!')
                model.stop_training = True
    callbacks = myCallback()

    # This Code Block should Define and Compile the Model. Please assume the images are 150 X 150 in your implementation.
    model = tf.keras.models.Sequential([
        tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3)),
        tf.keras.layers.MaxPool2D(2, 2),
        tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
        tf.keras.layers.MaxPool2D(2, 2),
        tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
        tf.keras.layers.MaxPool2D(2, 2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(40, activation='relu'),
        tf.keras.layers.Dense(1, 'sigmoid')
    ])

    model.compile(loss='binary_crossentropy', optimizer=RMSprop(lr=0.001), metrics=['accuracy'])

    # This code block should create an instance of an ImageDataGenerator called train_datagen
    # And a train_generator by calling train_datagen.flow_from_directory

    train_datagen = ImageDataGenerator(rescale=1./255)

    # Please use a target_size of 150 X 150.
    train_generator = train_datagen.flow_from_directory(
        'C:/Users/Alex Lai/Desktop-AJ0HRHM/Desktop/Deep Learning AI course/Week 4/test/happy-or-sad',
        target_size=(150, 150), batch_size=10, class_mode='binary'
    )
    # Expected output: 'Found 80 images belonging to 2 classes'

    # This code block should call model.fit_generator and train for
    # a number of epochs.
    # model fitting
    history = model.fit_generator(train_generator, steps_per_epoch=8, epochs=15, callbacks=[callbacks], verbose=1)
    # model fitting
    return history.history['accuracy'][-1]

# The Expected output: "Reached 99.9% accuracy so cancelling training!"
train_happy_sad_model()

```