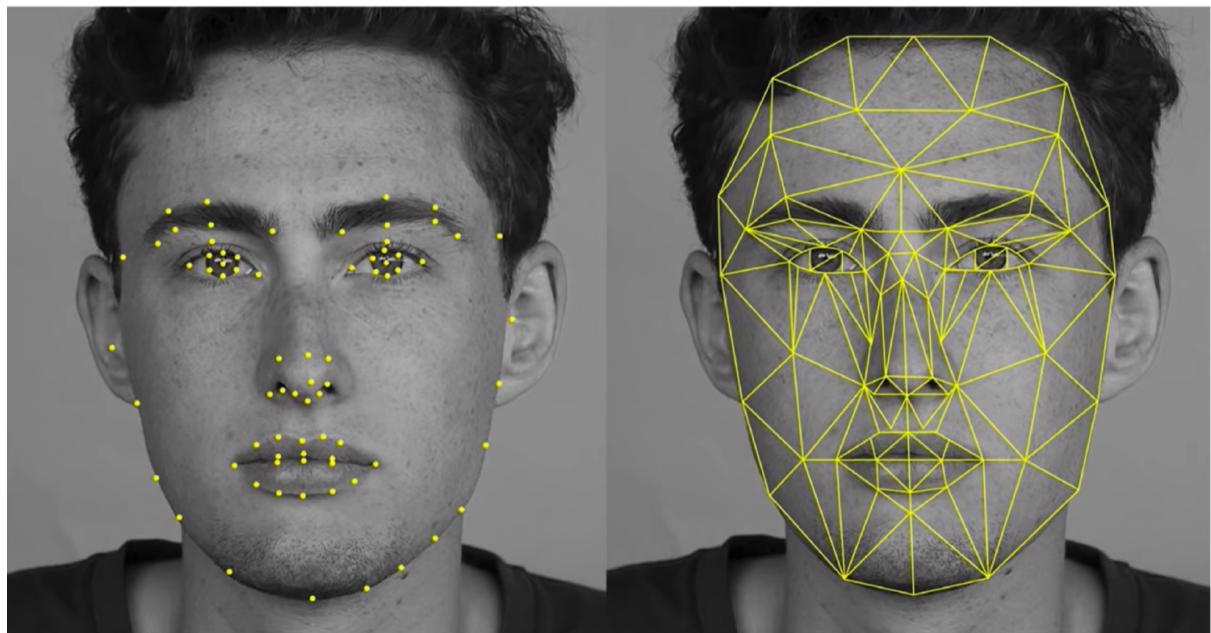




[3]Convolutional Neural Networks

Impacts

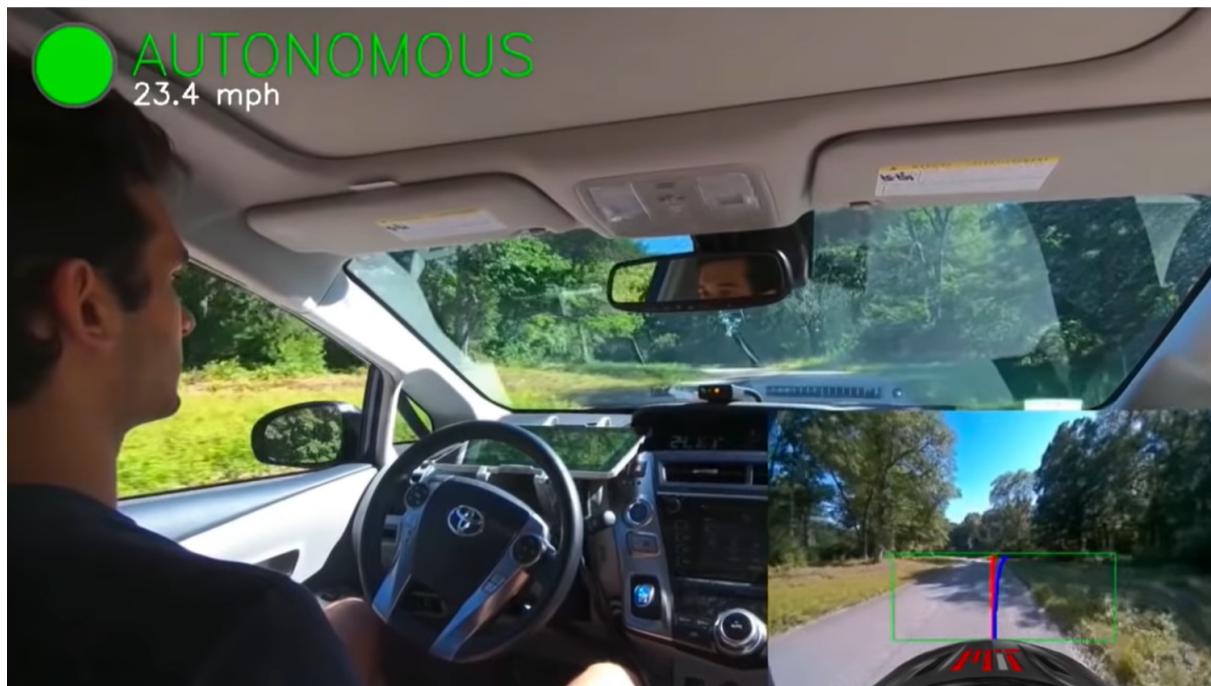
Facial Recognition



Medical, Biology, Healthcare

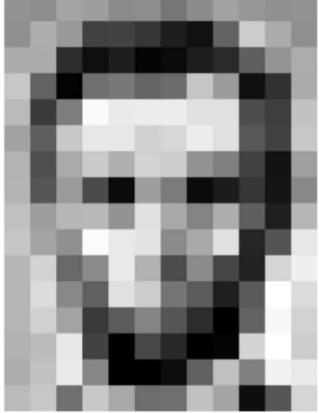


Self Driving Cars



What Computer "see"

- The computer sees the image as a two dimensional matrix with gray scale
- If the image is color (RGB) then the colors stack on to each other

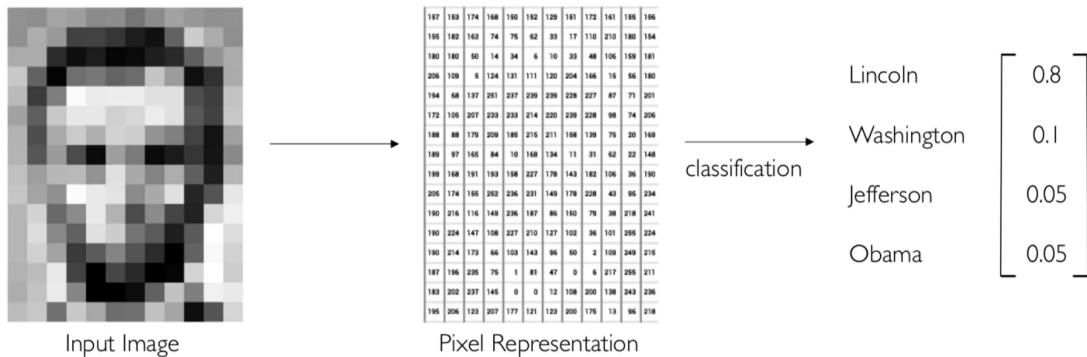


157	153	174	168	150	162	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	6	124	131	111	120	204	166	15	56	180
194	68	197	251	237	239	239	226	227	87	71	201
172	105	207	233	233	214	220	230	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

An image is just a matrix of numbers [0,255]!
i.e., 1080x1080x3 for an RGB image

Task for Computer Vision

- Regression and classification are the common tasks for computer vision
 - Regression: The output variable takes continuous value
 - Classification: output variable takes class label and can produce a probability of belonging to each class



High Level Feature Detection

- To tell the difference between each image, there are specific features of each image and that's a way to tell the difference between one image and another

Let's identify key features in each image category



Nose,
Eyes,
Mouth



Wheels,
License Plate,
Headlights



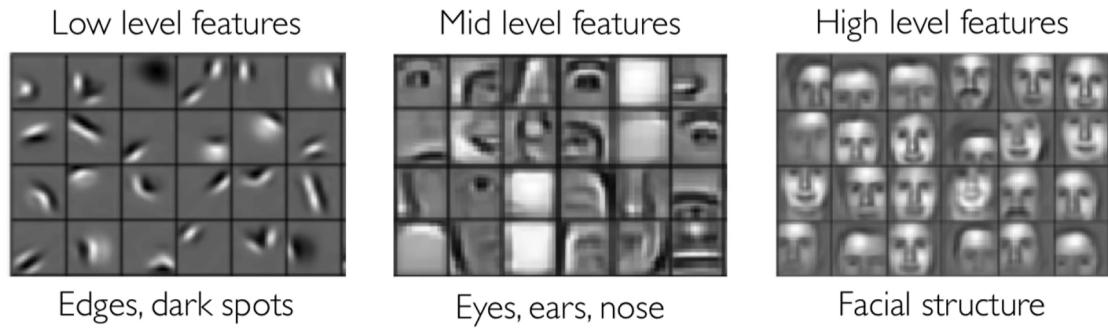
Door,
Windows,
Steps

Manual Feature Extraction

- The detection hierarchy is the bottleneck
- There could be lots of variation between images

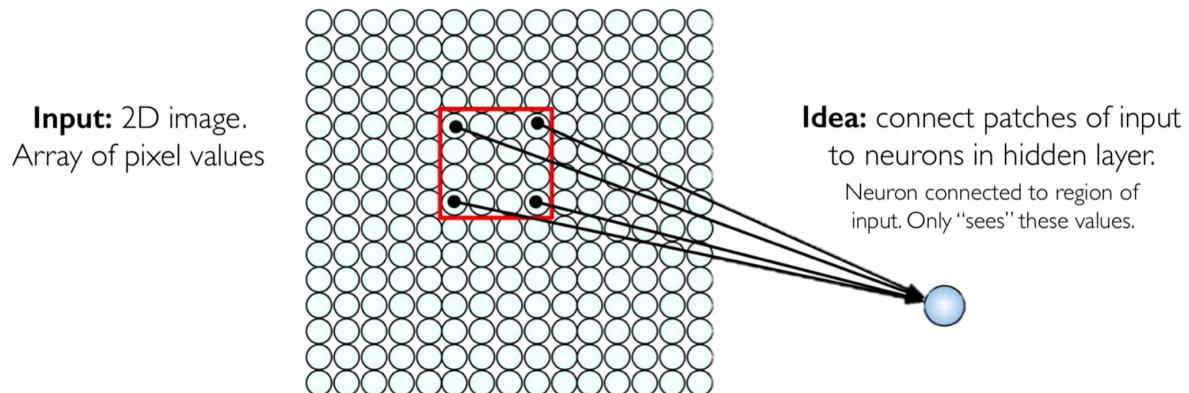
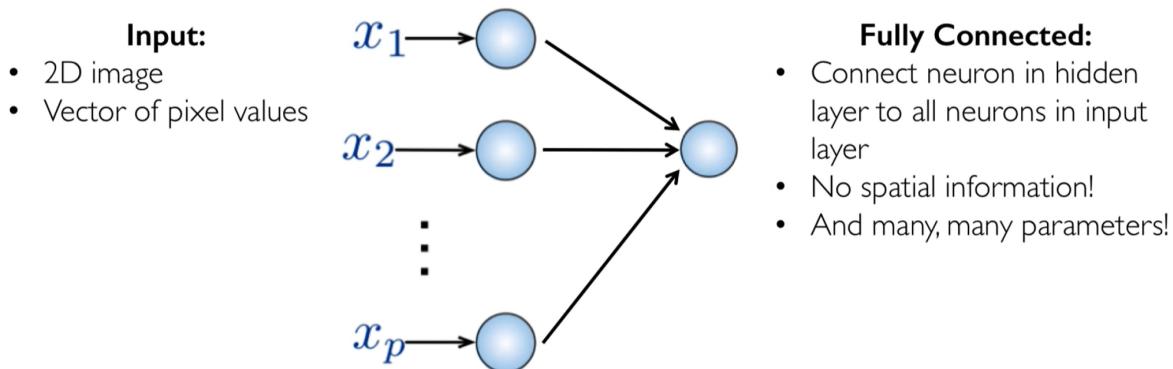


- We can learn a hierarchy of feature directly from the data instead of hand engineering

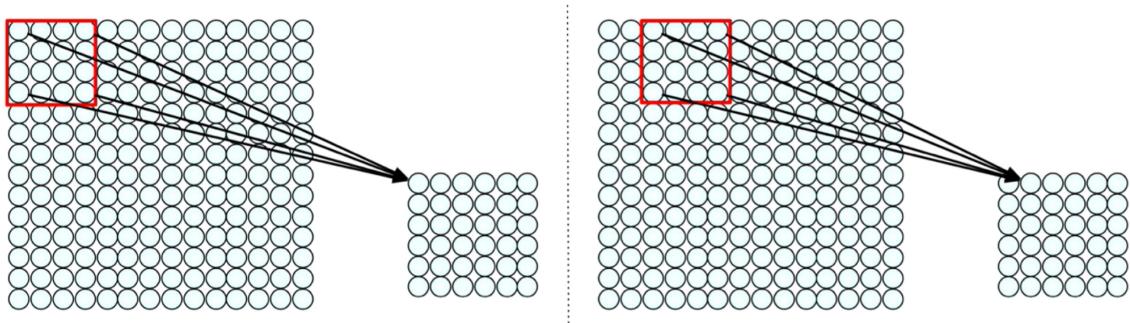


Learning Visual Features

- Using 2d image to compress that down into a 1d vector and putting it into the fully connected (Dense) layer



- Connect patches to input layers to a single neuron in subsequent layer
- Use a slide window to define the connection as shown below

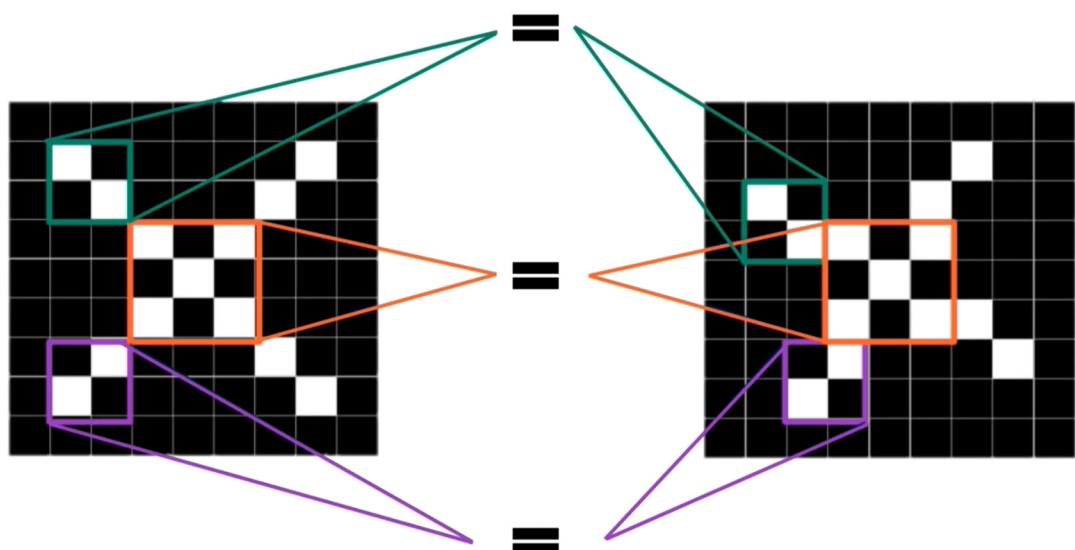


Feature Extraction with Convolution

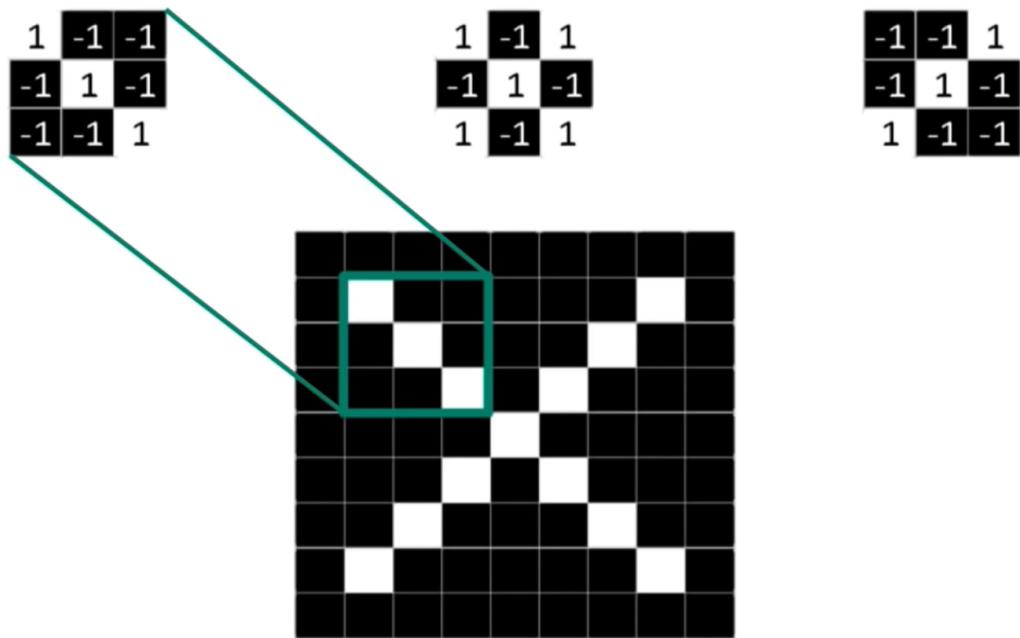
- Apply a set of weight, a filter to extract the local features
- Use multiple filters to extract different features
- spatially share parameters of each filter

Feature Extraction and Convolution

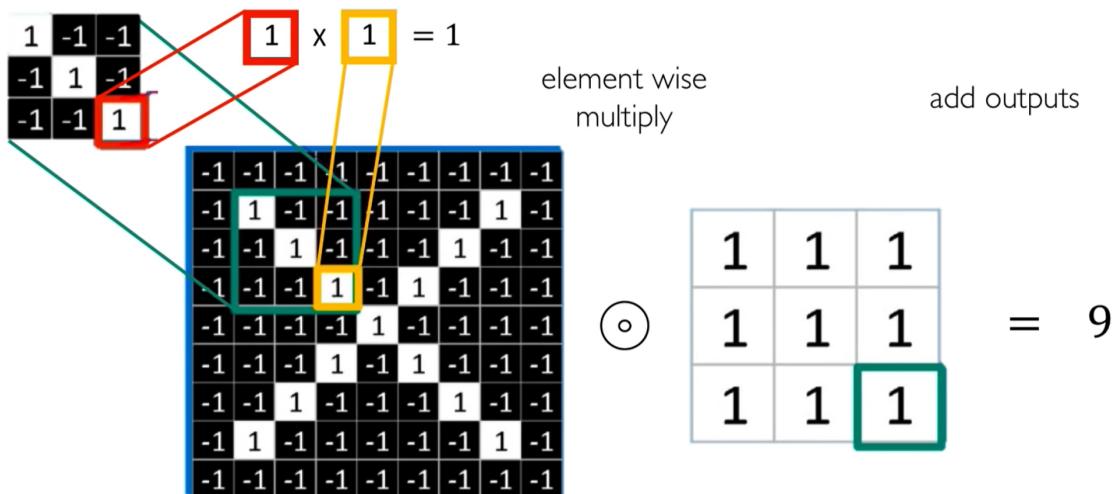
- Example below compares the weather the two x are the same



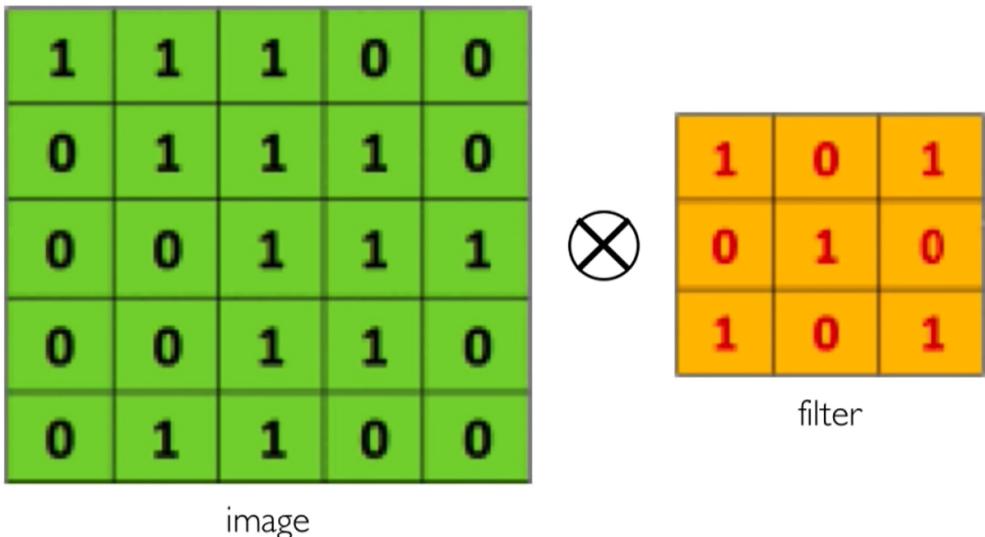
- The model could see that there are a few main features that are the same therefore they are both x
- Using filters, we want to capture features that is important to the image



- The image below shows what we do with the operation of the filters



- If we want to perform a convolution for a 5*5 grid with a 3*3 filter, the image shown below



- We want to slide the filter through the image and multiply the output
- Here are the steps of how it's done:

$$\begin{array}{c}
 \begin{matrix} 1_{x_0} & 1_{x_0} & 1_{x_0} & 0 & 0 \\ 0_{x_0} & 1_{x_1} & 1_{x_0} & 1 & 0 \\ 0_{x_0} & 0_{x_0} & 1_{x_0} & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} \otimes \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} = \begin{matrix} 4 & & \\ & & \\ & & \end{matrix} \\
 \text{filter} \qquad \qquad \qquad \text{feature map}
 \end{array}$$

$$\begin{array}{c}
 \begin{matrix} 1 & 1_{x_0} & 1_{x_0} & 0_{x_1} & 0 \\ 0 & 1_{x_0} & 1_{x_1} & 1_{x_0} & 0 \\ 0 & 0_{x_1} & 1_{x_0} & 1_{x_1} & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} \otimes \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} = \begin{matrix} 4 & 3 & \\ & & \end{matrix} \\
 \text{filter} \qquad \qquad \qquad \text{feature map}
 \end{array}$$

$$\begin{array}{c}
 \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1_{x_1} & 1_{x_0} & 1_{x_1} \\ 0 & 0 & 1_{x_0} & 1_{x_1} & 0_{x_0} \\ 0 & 1 & 1_{x_1} & 0_{x_0} & 0_{x_1} \end{matrix} \otimes \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} = \begin{matrix} 4 & 3 & 4 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{matrix} \\
 \text{filter} \qquad \qquad \qquad \text{feature map}
 \end{array}$$

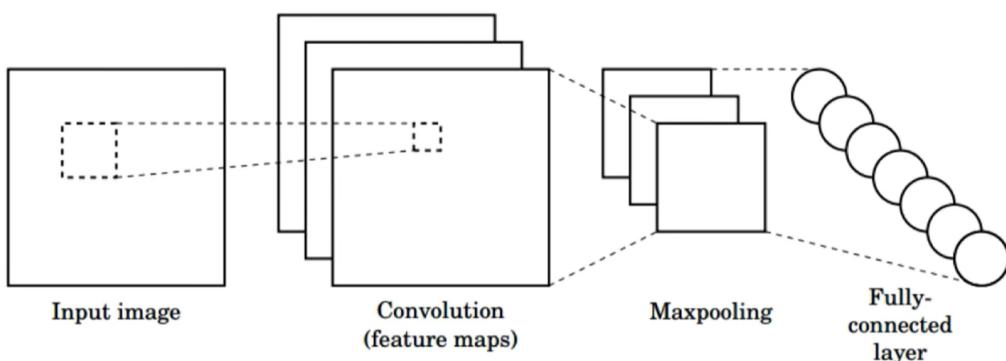
- Keep repeating this process until the feature map has been filled completely

Producing Feature Map

- By changing the filter of the image it could change what the convolution is looking for



Convolution Neural Networks (CNNs)



1. Convolution: Apply Filter to generate feature maps

`tf.keras.layers.Conv2D`

1. Non-linearity: Often ReLU

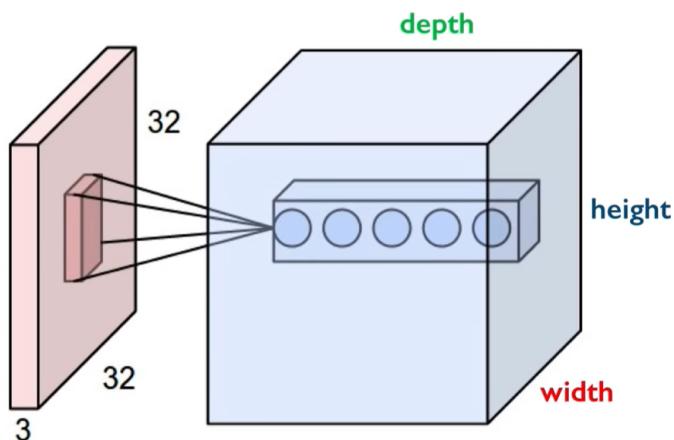
`tf.keras.layers.MaxPool2D`

1. Pooling: Down-Sampling operation of each feature map

`tf.keras.activations.*`

CNNs special arrangement of output volume

- The convolution output could be in multiple dimensions



Layer Dimensions:

$h \times w \times d$

where h and w are spatial dimensions
d (depth) = number of filters

Stride:

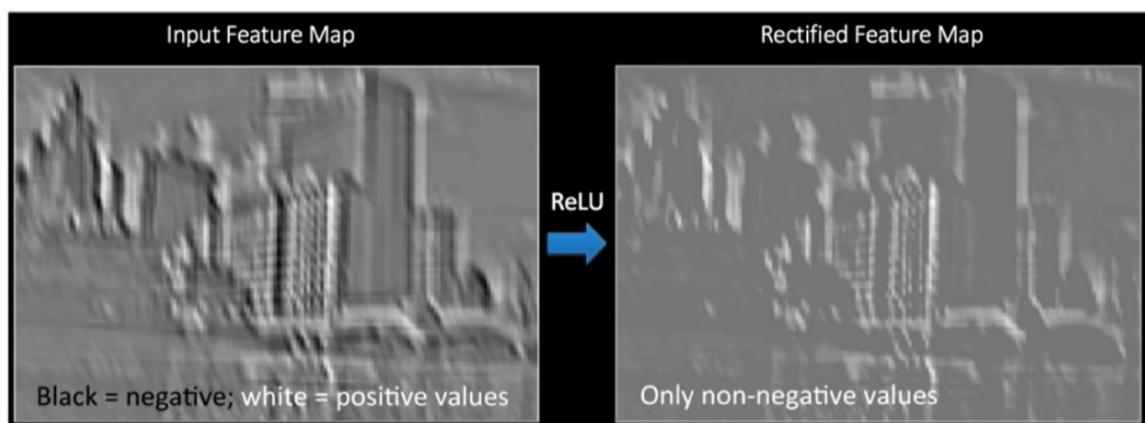
Filter step size

Receptive Field:

Locations in input image that
a node is path connected to

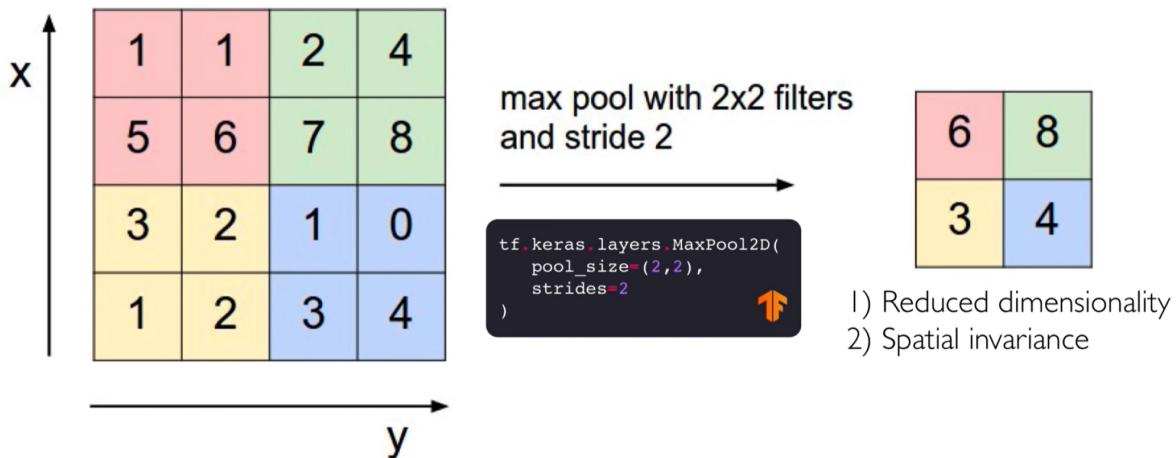
Non-linearity

- Non-linearity is applied after every convolution operation
- ReLU replaces all negatives to zero



Polling

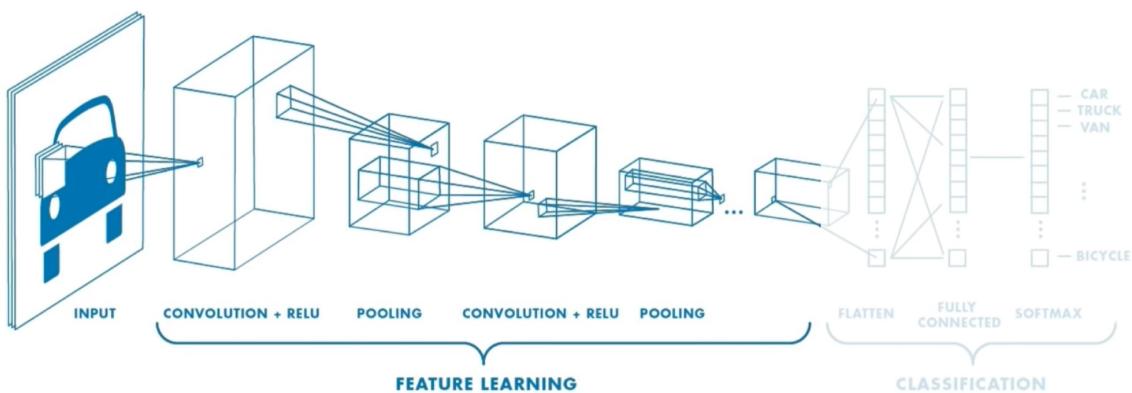
- Polling selects a greatest or lowest value in the specific area and replace the entire area with that single number



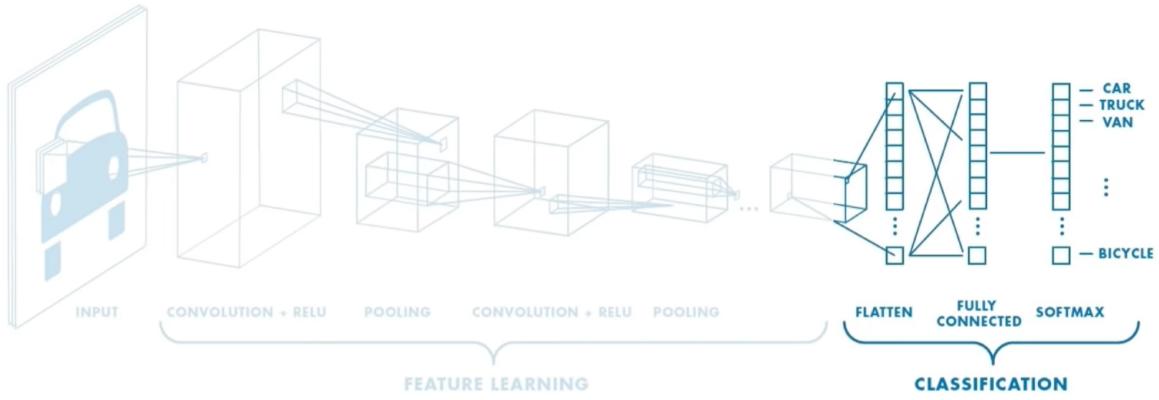
- This down-samples the image
- It allows us to shrink the image and still maintain the spatial structure

Steps for CNN networks

- The first part is Feature learning
 - Learns the image through convolution
 - introduces non-linearity
 - Reduce size and dimension through pooling

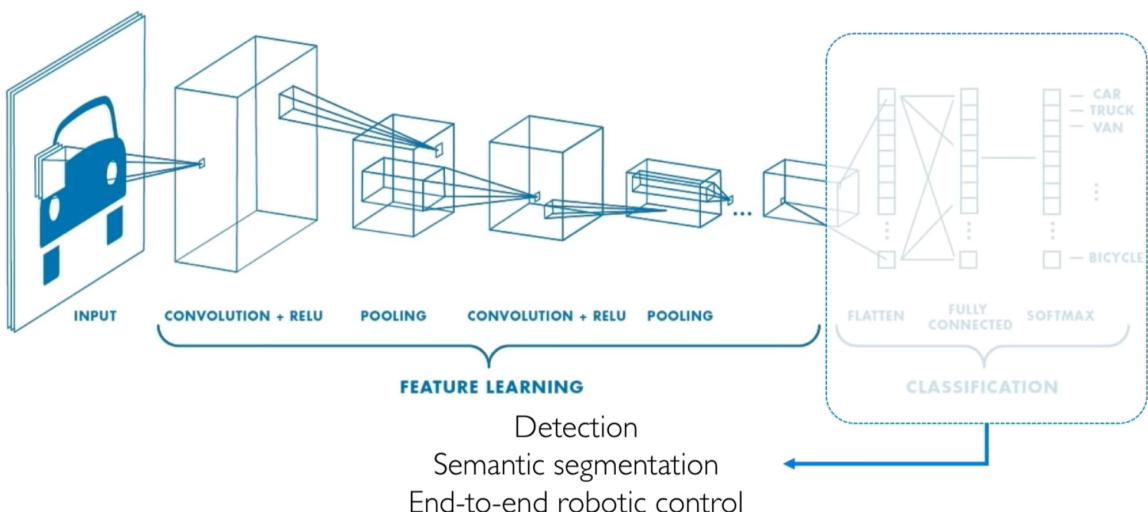


- The second part is classification
 - CONV and POOL output high level features of input
 - fully connected layers use this to classify the input image
 - express the probability of it belonging to a particular class
- The activation is softmax



Applications

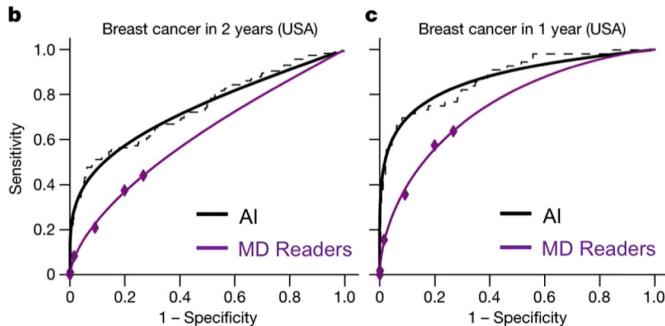
- CNNs not only can be used for classification, it could also be used for
 - Detection
 - Semantic segmentation
 - End-to-end robotic control



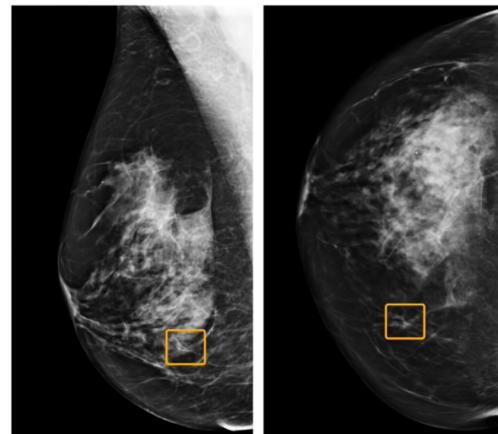
Detection: Breast Cancer Screening

International evaluation of an AI system for breast cancer screening

nature



CNN-based system outperformed expert radiologists at detecting breast cancer from mammograms



Breast cancer case missed by radiologist but detected by AI

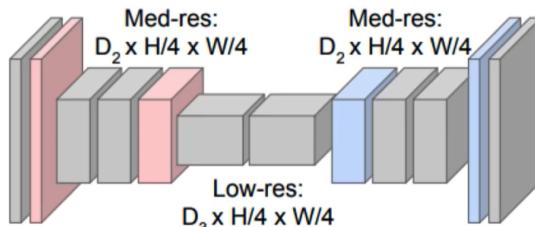
Semantic Segmentation: Fully Convolutional Network

FCN: Fully Convolutional Network.

Network designed with all convolutional layers, with **downsampling** and **upsampling** operations



Input:
 $3 \times H \times W$

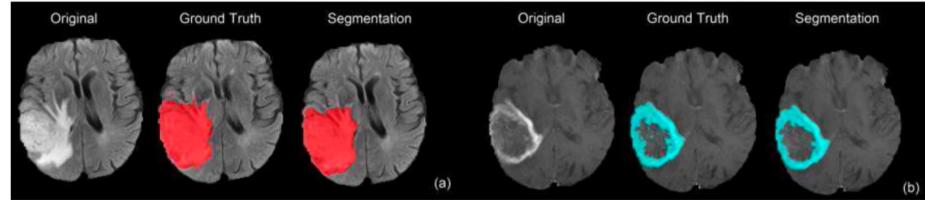


Predictions:
 $H \times W$

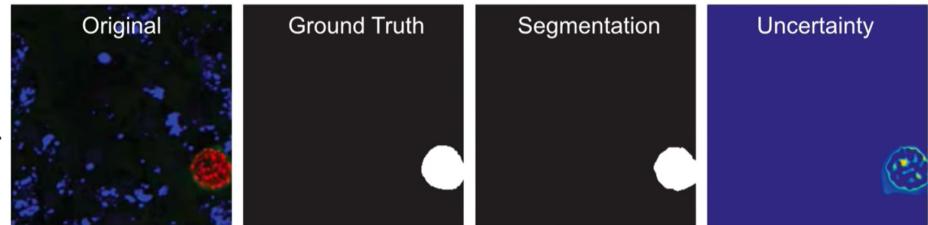
`tf.keras.layers.Conv2DTranspose`

Semantic Segmentation: Biomedical Image Analysis

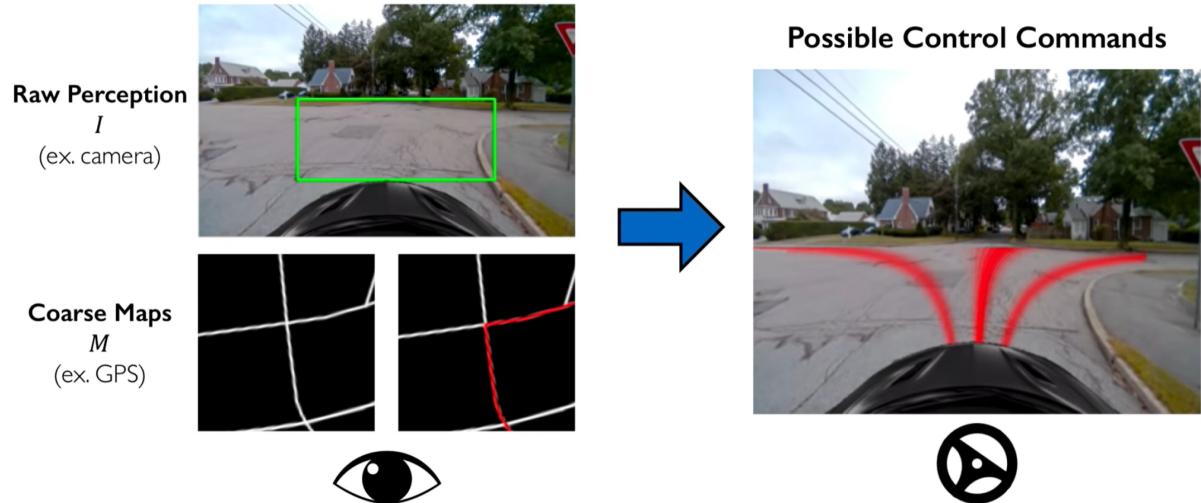
Brain Tumors
Dong+ MIUA 2017.



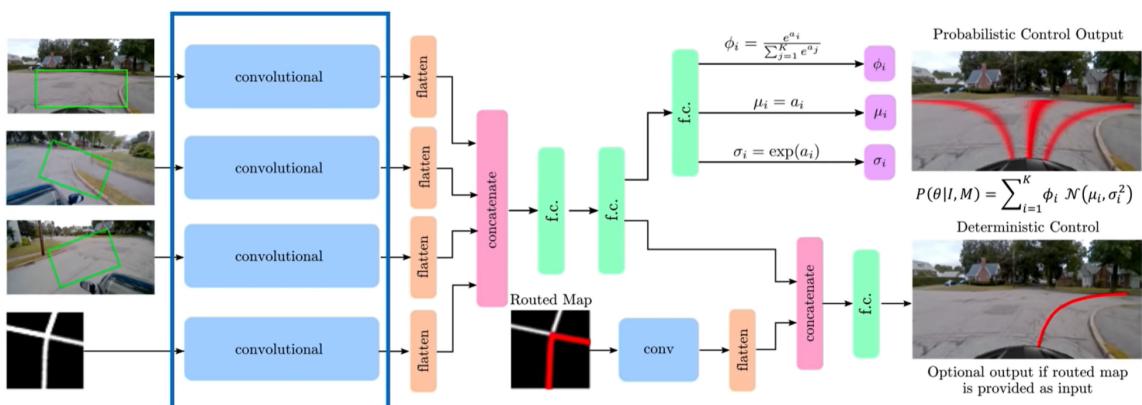
Malaria Infection
Soleimany+ arXiv 2019.



Self-Driving Car: Navigation from Visual Perception



Entire model is trained end-to-end **without any human labelling or annotations**



Summary

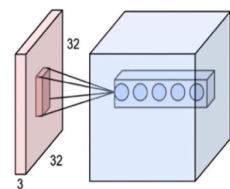
Foundations

- Why computer vision?
- Representing images
- Convolutions for feature extraction



CNNs

- CNN architecture
- Application to classification
- ImageNet



Applications

- Segmentation, image captioning, control
- Security, medicine, robotics

